

Epochs '26 – Assignment 11 is Live!

Welcome to the final challenge of **Epochs '26**!

In this assignment, you'll build a **PDF Question Answering Application** using **Retrieval-Augmented Generation (RAG)**. Your application should understand the contents of an uploaded PDF, answer user queries, and maintain conversation history for natural follow-up questions.

What You'll Do

- ◆ Load a PDF document.
 - ◆ Split the document into chunks and generate embeddings.
 - ◆ Store embeddings in a vector database for retrieval.
 - ◆ Integrate a **free LLM** to generate context-aware responses.
 - ◆ Implement **conversation memory** so the assistant can understand follow-up questions.
 - ◆ Build a simple and interactive user interface.
-

Suggested Technologies

- LangChain
- ChromaDB
- PyPDFLoader
- Sentence Transformers
- Gradio (*Recommended*)

Free LLM Options

-  Ollama
 -  Google Gemini API (Free Tier)
 -  Groq API
-

Submission

GitHub Repository

Include:

- Complete source code
- Project files
- `requirements.txt`

README.md

Include:

- Participant Name
- MUID
- Project Overview
- Technologies Used
- Memory Implementation
- Challenges Faced
- Future Improvements

Deployment Link (*Mandatory*)

Deploy your application using any suitable platform (e.g., Hugging Face Spaces, Render, Streamlit Community Cloud) and include the **public deployment URL**.

Submit **both your GitHub Repository link and Deployment link** in **#epochs-submissions** using the tag:

#evn-ds-epochs26-day11

★ **Reward: 300 Karma Points**

🕒 **Deadline: Within 48 hours.**